

Idea

1. search for all primes up to \sqrt{n}
 - reason (trivial): $a * b > n$ (if $a > b \vee b > a$) for $a, b, c, \in \mathbb{N}$
 - uses regular Sieve of Erathostenes
→ represented as *array* of 1 and 0
2. split array up into *segments*
 - instead of one giant *numbers* array, the array is only created with the length of a segment
 - orientate towards CPU *L1d cache* (`lscpu | grep "cache"`)
→ maximum efficiency
3. for each segment:
 - check for multiples of *all* primes found in 1)
4. in all of this: *Ignore* all even numbers
 - even numbers cannot be prime
 - allows for only checking for every *other* multiple of a prime p because:
 - for every $p \neq 2 \in \mathbb{P}$ is true: $2 \nmid p$
 - $2 \nmid p \rightarrow 2 \nmid (2a + 1)p \wedge 2 \mid 2a * p$ ($a \in \mathbb{N}_0$)
 - since $2p$ is obviously *even*, it gets ignored by *every* part of the program, thus its state is irrelevant
 - ⇒ this means that when counting primes, *2 is never included* → thus, the number of found primes is always *one less* than the actual number of primes

```
is_prime[] = regular_sieve(until sqrt(n))
num_primes = 1 # compensates for ignoring 2
for each segment [low, high]:
    sieving_primes = [primes in is_prime until sqrt(high)]
    multiples: list[len(sieving_primes)]

    for each prime p in sieving_primes:
        a = multiples[index of p] # a is always uneven
        for each multiple m = 2a * p:
            eliminate m from segment
        when multiple > high:
            store m in multiples

    num_primes += number of primes in segment
```

Why it is faster

Locality of Memory

A regular sieve is *massively* inefficient when it comes to caching, as each time, the processor has to go over the *entire* array. Splitting it up into smaller chunks eases the workload of the processor, and also, after finishing each segment, all non-primes disappear from the cache. Making it for once more space-efficient, but also more time-efficient due to *locality of memory* - whatever that is exactly.

Ignoring even numbers

Only dealing with uneven numbers essentially *halves* the workload. This can also be observed in the time: When processing *all* numbers, runtime jumps from *~250 ms* all the

Implementation:

[Implementation Segmented Sieve](#)