

# Implementation Segmented Sieve

## Code

[segmented\\_sieve.c](#)

## Runtime

**Time Complexity:**  $O(n \log(n))$

**Space Complexity:**  $O(\sqrt{n})$

**Timed** (average-case):

as *linux executable*: ~225 ms

as *windows executable* under `wine`: ~260 ms

## Comparison

This compares the *average case* approximate runtime of the following entities:

- Segmented Sieve (with hardcoded values): [segmented\\_sieve\\_hardcoded.c](#)
- Segmented Sieve: [segmented\\_sieve.c](#)
  - under `wine`: Compiled to `exe`, then ran via the wine linux library. This serves to better compare the Benchmark, which is provided as a windows executable.
- Benchmark: [https://elearning.dhbw-stuttgart.de/moodle/pluginfile.php/660507/mod\\_folder/content/0/PrimDynC.exe?forcedownload=1](https://elearning.dhbw-stuttgart.de/moodle/pluginfile.php/660507/mod_folder/content/0/PrimDynC.exe?forcedownload=1) --> *PrimDynC.exe*
  - This is the program that also allows for User Input. *PrimOpt2*, which does not do so, is faster, but is hardcoded to  $10^8$ .

**$10^8$**

	Segmented Sieve (with hardcoded values)	Segmented Sieve	Segmented Sieve (under wine)	Benchmark (under wine)	Benchmark (with hardcoded values)
battery power	230 ms	280 ms	340 ms	850 ms	700 ms
performance mode	225 ms	?? ms	?? ms	?? ms	340 ms

--> As *space complexity* is optimized, Segmented Sieve runs basically equally fast regardless of available resources

--> `wine` emulator adds 20% of runtime

10<sup>9</sup>

	Segmented Sieve (with hardcoded values)	Segmented Sieve	Segmented Sieve (under wine)	Benchmark (under wine)
battery power	2450 ms	2900 ms	3300 ms	6900 ms

10<sup>10</sup>

Sadly, under `wine`, both programs don't run properly anymore, as they're lacking storage space. This is a shame.

The regular [segmented\\_sieve.c](#) needs ~ 35 seconds. The regular [segmented\\_sieve.c](#) needs ~35 seconds.

## Documentation

since I need `stdio` for this and it redefines `timeval` from `linux/time.h`, I'm not working with the provided Nimmzeit module, instead using `clock()` to measure the time. A giant Fuck You to whoever wrote that stupid module.

Futhermore, since the evaluation happens *inside* the sieving, it cannot be timed separately. Of course one could time *each* individual sub-evaluation, however that slows the program down by ~50 ms, which is about 20% of total runtime and thus not feasible.

## Storage of Sieving primes

### Structure `list_uint64`

Roughly emulates a *list* in python.

```
typedef struct
{
    uint64_t *data;
    uint64_t max_size;
    uint64_t length;
} list_uint64;
```

### Attributes

- `data` -> Pointer to where the *actual data* of the list lies. This can be used synonymously with an array using bracket syntax: `list.data[index]`.
- `max_size` -> How much *memory* has been allocated. Is used when creating the `multiples` list, to keep its size in sync with the `primes` list.

- *length* -> equivalent to the return value of `len(list)` in python. Stores how many elements the list currently holds. Can be used to iterate over only the parts of the list that actually store values.

## Initialization with default values

*Parameter size*: The maximum size the List should be.

*Return value*: The newly created list. Since structs don't need explicit pointers to be passed around, the list instance can simply be returned.

```
// creates a new list filled with default values
list_uint64 new_list(uint64_t size)
{
    list_uint64 list = {
        malloc(size * sizeof(uint64_t)),
        size,
        0
    };

    return list;
}
```

- *data*: Manually allocates  $size * sizeof(uint64_t)$  of memory. This way there can be *size* Elements of the type *uint64\_t* stored in the list.
- *max\_length*: This is initialized to the given size of the wanted list.
- *length*: Trivial - As there are no initial elements in the list, its size should be 0.

## Appending

```
// appends an item to given list. List is passed as reference.
void list_append(list_uint64 *list, uint64_t value) {
    list->data[list->length++] = value;
}
```

This is more or less a very fancy *array syntax*. It first accesses the data list via pointer operations (`list->data`), then stores the value at the end (`list->length`) and then increases the length attribute by one.

## Setup

1. create `is_prime` array to later run a regular sieve on
2. fill this array with one's --> leave indexes for 0 and 1 out, because those aren't primes
3. run a regular [Sieve of Erathostenes > C implementation](#) on the `is_prime` array
  - > places all indexes of primes until  $\sqrt{n}$  with 1, else 0
  - > using `is_prime[index]` will return whether *index* is *prime*

4. create a `list_uint64` instance to later store the sieving primes in
  - The size of the array is approximated: [Approximation of Primes](#)

```
bool is_prime[(size_t)sqrt(upper_limit)];
memset(is_prime + 2, true, sizeof(is_prime) - 2);
sieve(is_prime, sizeof(is_prime));

list_uint64 primes = new_list(PRIMES_UNTIL_SQRT);
```

## Approximation of Primes

The [prime number theorem](#) suggests that  $\lim_{x \rightarrow \infty} x \log_{10}(x) = \pi(x)$

Since  $x \log(x)$  generally falls short of the actual number of primes, however we want to rather *overshoot* than *underestimate*, The result has to be increased. by some amount.

Let  $A_x = \frac{\pi(x)}{x \log_{10}(x)}$

x	$\pi(x)$	$A_x$
10	4	0.921
$10^2$	25	1.151
$10^3$	168	1.161
$10^4$	1229	1.132
$10^5$	9592	1.104
$10^6$	78498	1.084
$10^7$	664579	1.071
$10^8$	5761455	1.061
$10^9$	50847534	1.054
$10^{10}$	455052511	1.048

In this table can be seen, that  $\forall_x A_x \leq 1.52$ . So, to get a number a bit larger than the actual number of primes until a parameter  $n$ , this code can be used:

```
int n;
(n / (uint64_t)log10(n)) * 1.2
```

## Algorithm

### 1. Initialization of relevant variables

#### Scope of function

Variables that will be used over and over again and incremented in each segment

```
uint64_t low, high;
uint64_t num_primes = 1;
uint64_t sieving_prime_candidate = 3;
uint64_t prime_counter = 3;
list_uint64 multiples = new_list(sieving_primes.max_size);

bool segment[segment_size];
```

- *low, high*: control boundaries of each segment
- *num\_primes*: serves as counter for the total number of primes
  - each segment adds the number of primes it contains to this variable
  - variable is initialized to 1 because the entire code ignores multiples of 2
- *sieving\_prime\_candidate*: bit of a chunky name, but tells exactly what the variable is for
  - value will increase with each segment
  - this variable will get checked against `is_prime` to see if its value is prime
- *prime\_counter*: used at the end of each segment to iterate through the `segment` array to find all primes
  - not to be confused with *num\_primes*
- *multiples*: stores the first multiple of a prime that is *outside* the current segment, so the next segment will know where to start eliminating for the prime
  - since there is exactly *one* "next multiple" for every sieving prime, the lists have the same size, and `multiples.data[i]` will always correspond to `sieving_primes.data[i]`
- *segment*: the good old essential Sieve of Eratosthenes array for a given segment.
  - will be reused for different segments
  - --> data of one segment *does not persist*
  - attention: since the entire program *ignores even numbers*, every even index will also be marked as *true*, just not evaluated.

## Scope of outer loop

Sets environment for the current segment.

```
for (low = 0; low <= upper_limit; low += segment_size)
{
    memset(segment, true, sizeof(segment));
    high = low + segment_size - 1;
    if (high > upper_limit)
        high = upper_limit;
```

- Loop condition is trivial

- sets every element of the `segment` array to *true*
- high is either set to *one less* than the low of the next segment or the upper limit (trivial)

## 2. Filtering of Sieving Primes

*Sieving Prime*: A prime  $p$  that is used to cross of multiples  $a * p$  within a segment  $[l, n]$ .  
 Sieving Primes only need to be smaller than  $\sqrt{n}$ .

```
for (; sieving_prime_candidate * sieving_prime_candidate <= high;
    sieving_prime_candidate += 2)
{
    if (is_prime[sieving_prime_candidate])
    {
        list_append(&sieving_primes, sieving_prime_candidate);
        list_append(
            &multiples,
            (sieving_prime_candidate * sieving_prime_candidate)
- low
        );
    }
}
```

- sieving prime candidates always get increased by 2  
 --> since `sieving_prime_candidate` is initialized to 3, this ignores *all even numbers*
- If a candidate is in fact a prime, it is appended to the *list* `sieving_primes`
  - also, a corresponding entry is appended to the `multiples` list to be able to be used *later on*
  - The default value for the new `multiples` entry can be derived from the common standard sieve optimization, to start eliminating from  $p * p$  instead of  $p$

## 3. Eliminating multiples

```
for (uint64_t i = 0; i < sieving_primes.length; i++)
{
    uint64_t j = multiples.data[i];
    for (; j < segment_size; j += sieving_primes.data[i] * 2)
    {
        segment[j] = false;
    }
    multiples.data[i] = j - segment_size;
}
```

- Iterator variable  $j$  is assigned the corresponding `multiples` entry  
 --> if the sieving prime  $p$  has been used in the segment before, this will be the first multiple of  $p$  in the current segment

- Since all even indexes get ignored, we can only check every *other* multiple of  $p$ 
  - $2 \nmid p \rightarrow 2 \nmid (2a + 1)p \wedge 2 \mid 2a * p$ , therefore  $2a * p \notin \mathbb{P}$
- Finally, the `multiples` entry is updated to store the *next* multiple, that doesn't lie in the current segment
  - this usage of  $j$  is why  $j$  needed to be declared *before* the for loop

## 4. Evaluation

Since data of an individual segment doesn't persist, the evaluation needs to take place inside of a segment.

```
for (; prime_counter <= high; prime_counter += 2)
{
    if (segment[prime_counter - low])
    {
        num_primes++;
    }
}
```

- Works the same way as the evaluation of a regular Sieve of Eratosthenes.

## Idea

### Segmented Sieve

1. search for all primes up to  $\sqrt{n}$ 
  - reason (trivial):  $a * b > n$  (if  $a > \sqrt{n} \wedge b > \sqrt{n}$ ) for  $a, b, c, \in \mathbb{N}$
  - uses *regular* [Sieve of Eratosthenes](#)  
--> represented as *array* of 1 and 0
2. split array up into *segments*
  - instead of one giant *numbers* array, the array is only created with the length of a segment
3. for each segment:
  - check for multiples of *all* primes found in 1.
4. in all of this: *Ignore* all even numbers
  - even numbers cannot be prime
  - allows for only checking for every *other* multiple of a prime  $p$  because:
    - for every  $p \neq 2 \in \mathbb{P}$  is true:  $2 \nmid p$
    - $2 \nmid p \rightarrow 2 \nmid (2a + 1)p \wedge 2 \mid 2a * p$  ( $a \in \mathbb{N}_0$ )
    - > since  $2a * p$  is obviously *even*, it gets ignored by every part of the program, thus its state is irrelevant

=> this means that when counting primes, *2 is never included* --> thus, the number of found primes is always *one less* than the actual number of primes

## pseudo code

```
is_prime[] = regular_sieve(until sqrt(n))
num_primes = 1 # compensates for ignoring 2
for each segment [low, high]:
    sieving_primes = [primes in is_prime until sqrt(high)]
    multiples: list[len(sieving_primes)]

    for each prime p in sieving_primes:
        a = multiples[index of p] # a is always uneven
        for each multiple m = 2a * p:
            eliminate m from segment
        when multiple > high:
            store m in multiples

    num_primes += number of primes in segment
```

## Why it is faster

### Locality of Memory

A regular sieve is *massively* inefficient when it comes to caching, as each time, the processor has to go over the *entire* array. Splitting it up into smaller chunks eases the workload of the processor, and also, after finishing each segment, all non-primes disappear from the cache. Making it for once more space-efficient, but also more time-efficient due to *locality of memory* - whatever that is exactly.

### Ignoring even numbers

Only dealing with uneven numbers essentially *halves* the workload. This can also be observed in the time: When processing *all* numbers, runtime jumps from *~250 ms* all the way to *~500 ms*

## Implementation

[Implementation Segmented Sieve](#)

## Reference



explanation of algorithm: <https://github.com/kimwalisch/primesieve/wiki/Segmented-sieve-of-Eratosthenes>